

F03 | Completed repairs versus introduced errors

Repair Without Rewrite / DEVELOPMENT evidence / seed 42

Failure-inclusive natural counts; no stacking of rates with different denominators

Prescribed endpoint	Completed repair / 1,413 errors	Raw repair / 1,413 errors	Introduced / 50,926 words	Output errors / 50,926 words	Invalid/capped / 2,696 cases
RAW (identity)	0	0	0	1,413	0
B100 D1-U8	44-45	44-45	15,628-15,629	16,997	0
C101 D1-U8	18	20	2,702	4,095	26
ByT5 ten-pass	120	120	357	1,650	0

RAW repair and introduction are zero by identity. Alignment ranges are identification bounds, not confidence intervals.

Raw conservation: output errors = RAW errors - raw repair + introduced errors.

C101 conservation uses 20 raw repairs; its 18 completion-qualified repairs are a distinct utility measure.

Originally lexical-zero stratum (retrospective calculation; limited reconstruction)

Endpoint	Damaged cases / 1,818	Errors in that stratum (count)
B100 D1-U8		1,60010,151
C101 D1-U8		2881,614
ByT5 ten-pass		183239

Historical endpoint scores were independently reconstructed; detailed retrospective stratum decomposition was not fully requalified in frontier v3. No new subgroup extraction is performed.

Training-seed and domain/pretraining uncertainty remain unmeasured; all prescribed endpoints failed WER below RAW.