

F07 | Whole-output benefit and reference-aware oracle headroom

Repair Without Rewrite / DEVELOPMENT evidence / seed 42

Fixed complete proposals: outcome categories on all natural DEVELOPMENT cases

Fixed endpoint	Better /2,696	Equal /2,696	Worse /2,696	Actual errors /50,926 words
B100 D1-U8	2	297	2,397	16,997
C101 D1-U8	3	2,197	496	4,095
ByT5 exact ten-pass	67	2,360	269	1,650

Reference-aware oracle: NON-DEPLOYABLE ceiling (whole-output choice only)

Fixed endpoint	Beneficial complete cases /2,696	Beneficial groups /52	Removable errors /1,413	Oracle errors /50,926 words	Oracle WER % non-deployable
B100 D1-U8	2	2	2	1,411	2.770687
C101 D1-U8	3	3	3	1,410	2.768723
ByT5 exact ten-pass	67	36	109	1,304	2.560578

RAW baseline: 1,413 errors / 50,926 words; WER % = **2.774614**

Recorded oracle-gain sampling intervals; partial independent reconstruction

Non-deployable ceiling	95% descriptive group interval (WER pp)
B100 reference-aware oracle	[0.000000, 0.010418]
C101 reference-aware oracle	[0.000000, 0.013181]
ByT5 reference-aware oracle	[0.154027, 0.288196]

Oracle chooses RAW or one fixed complete proposal using reference errors; incomplete proposals choose RAW.

No edit mixing, new text, checkpoint choice or source-only policy is demonstrated.

Point gains were independently reaggregated in frontier v3; detailed feasibility intervals were not fully repeated.

Group sampling is not training-seed uncertainty and does not predict selectable benefit or transfer.