

F08 | Native training and evaluation runtime

Repair Without Rewrite / DEVELOPMENT evidence / seed 42

Physical recipe wall seconds: 13 attempts; components below are already included

Generation / recipe	Recipe wall seconds MEASURED	Completed updates	Unique presentations
G1 B100-lr1e-04	2117.458817	305	134,591
G1 C101-lr1e-04	5360.392531	305	134,591
G1 B100-lr3e-04	2255.100029	305	134,591
G1 C101-lr3e-04	5544.691853	305	134,591
G1 B100-lr6e-04	2468.490537	305	134,591
G1 C101-lr6e-04	5382.778655	305	134,591
G2-B100-D0-U8	10286.654194	2,440	134,591
G2-C101-D0-U8	9010.581726	2,440	134,591
G2-B100-D1-U1	7276.738185	305	136,755
G2-C101-D1-U1	8864.873318	305	136,755
G2-B100-D1-U8	8631.517741	2,440	136,755
G2-C101-D1-U8	8189.656452	2,440	136,755
G2-ByT5-D1-10pass	79862.099614	35,283	141,130

G1: 396-case evaluations; G2: 2,984-case evaluations. These are unequal-work native timings.
Archived D0-U1 controls are rescored, not retrained; they incur zero new G2 training.

ByT5 components: nested inside its recipe interval; do not sum into recipe time

Observation state	Actual presentations	Observation wall s in recipe wall	Decode s in observation wall	Scorer s in observation wall
Initialization	0	8525.970953	7901.264797	5.389383
2-pass nominal	28,228	7214.049868	5343.791777	11.141651
5-pass nominal	70,568	10059.125263	7636.503629	14.241679
10-pass exact	141,130	7074.588477	4916.474747	10.757864

ByT5 logged training seconds (MEASURED; included in recipe wall): 21143.834539

G1 endpoint decoding on the complete 396-case panel

G1 endpoint	Decode seconds in recipe wall	Decoder positions	Native update host s in recipe wall
B100-lr1e-04	62.872234	13,145	1278.312249
C101-lr1e-04	14.056227	2,833	4934.404539
B100-lr3e-04	80.771583	18,572	1357.819926
C101-lr3e-04	12.402017	2,264	5077.106274
B100-lr6e-04	72.899339	16,661	1553.420007
C101-lr6e-04	11.978356	2,120	4972.247960

Timing scope and unavailable cost

Recipe wall includes training, saves and prescribed evaluation. Decode/scorer lie inside observation wall.
G1 native update host intervals and ByT5 logged training intervals are components, not extra charges.
No unnamed timing remainder is assigned to a new category. Failed CPU checks are separate from scientific replay.
Exact device kernel time: UNMEASURED. G2 all-in operational span: UNMEASURED.
Electricity, depreciation and local monetary cost: UNPRICED. Upstream pretraining cost: UNESTABLISHED.
One local native implementation; failed models and unequal comparator training do not establish overall efficiency.